

ECE595 / STAT598: Machine Learning I

Lecture 18 Multi-Layer Perceptron

Spring 2020

Stanley Chan

School of Electrical and Computer Engineering
Purdue University



Outline


Discriminative Approaches

- Lecture 16 Perceptron 1: Definition and Basic Concepts
- Lecture 17 Perceptron 2: Algorithm and Property
- **Lecture 18 Multi-Layer Perceptron: Back Propagation**

This lecture: Multi-Layer Perceptron: Back Propagation



- Multi-Layer Perceptron
 - Hidden Layer
 - Matrix Representation
- Back Propagation
 - Chain Rule
 - 4 Fundamental Equations
 - Algorithm
 - Interpretation

Single-Layer Perceptron



- Input neurons x
- Weights w
- Predicted label $= \sigma(w^T x + w_0)$.



Multi-Layer Network



<https://towardsdatascience.com/multi-layer-neural-networks-with-sigmoid-function-deep-learning-for-rookies-2-bf464f09eb7f>


- Introduce a layer of **hidden** neurons
- So now you have two sets of weights: from input to hidden, and from hidden to output

Many Hidden Layers



- You can introduce as many hidden layers as you want.
- Every time you add a hidden layer, you add a set of weights.

Understanding the Weights



- Each hidden neuron is an **output** of a perceptron
- So you will have

$$\begin{bmatrix} h_1^1 \\ h_2^1 \\ \vdots \\ h_n^1 \end{bmatrix} = \begin{bmatrix} w_{11}^1 & w_{12}^1 & \dots & w_{1n}^1 \\ w_{21}^1 & w_{22}^1 & \dots & w_{2n}^1 \\ \vdots & \vdots & \ddots & \vdots \\ w_{m1}^1 & w_{m2}^1 & \dots & w_{mn}^1 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_m \end{bmatrix}$$

Progression to DEEP (Linear) Neural Networks

- Single-layer:

$$h = \mathbf{w}^T \mathbf{x}$$

- Hidden-layer:

$$\mathbf{h} = \mathbf{W}^T \mathbf{x}$$

- Two Hidden Layers:

$$\mathbf{h} = \mathbf{W}_2^T \mathbf{W}_1^T \mathbf{x}$$


- Three Hidden Layers:

$$\mathbf{h} = \mathbf{W}_3^T \mathbf{W}_2^T \mathbf{W}_1^T \mathbf{x}$$


- A LOT of Hidden Layers:

$$\mathbf{h} = \mathbf{W}_N^T \dots \mathbf{W}_2^T \mathbf{W}_1^T \mathbf{x}$$

Interpreting the Hidden Layer




A Trained Neural Network in which each hidden neuron classifies certain pattern



Feeding a handwritten digit of 0 should trigger the 4 hidden layer neurons, and then the first output neuron

- Each hidden neuron is responsible for certain features.
- Given an object, the network identifies the most likely features.

Interpreting the Hidden Layer



Two Questions about Multi-Layer Network

- How do we **efficiently** learn the weights?
 - Ultimately we need to minimize the loss

$$J(\mathbf{W}_1, \dots, \mathbf{W}_L) = \sum_{i=1}^N \|\mathbf{W}_L^T \dots \mathbf{W}_2^T \mathbf{W}_1^T \mathbf{x}_i - \mathbf{y}_i\|^2$$

- One layer: Gradient descent. Multi-layer: Also gradient descent, also known as Back propagation (BP) by Rumelhart, Hinton and Williams (1986)
- Back propagation = Very careful book-keeping and chain rule
- What is the optimization **landscape**?
 - Convex? Global minimum? Saddle point?
 - Two-layer case is proved by Baldi and Hornik (1989)
 - All local minima are global.
 - A critical point is either a saddle point or global minimum.
 - L -layer case is proved by Kawaguchi (2016). Also proved L -layer nonlinear network (with sigmoid between adjacent layers.)

Outline

Discriminative Approaches

- Lecture 16 Perceptron 1: Definition and Basic Concepts
- Lecture 17 Perceptron 2: Algorithm and Property
- **Lecture 18 Multi-Layer Perceptron: Back Propagation**

This lecture: Multi-Layer Perceptron: Back Propagation

- Multi-Layer Perceptron
 - Hidden Layer
 - Matrix Representation
- Back Propagation
 - Chain Rule
 - 4 Fundamental Equations
 - Algorithm
 - Interpretation

Back Propagation: A 20-Minute Tour

- You will be able to find **A LOT OF** blogs on the internet discussing how back propagation is being implemented.
- Some are **mystifying** back propagation
- Some literally just teach you the procedure of back propagation without telling you the intuition
- I find the following online book by Mike Nielsen fairly well-written
- <http://neuralnetworksanddeeplearning.com/>
- The following slides are written based on Nielsen's book
- We will not go into great details
- The purpose to get you exposed to the idea, and de-mystify back propagation
- As stated before, back propagation is chain rule + very careful book keeping

Back Propagation

- Here is the loss function you want to minimize:

$$J(\mathbf{W}_1, \dots, \mathbf{W}_L) = \sum_{i=1}^N \|\sigma(\mathbf{W}_L^T \dots \sigma(\mathbf{W}_2^T \sigma(\mathbf{W}_1^T \mathbf{x}_i))) - \mathbf{y}_i\|^2$$

- You have a set of nonlinear activation functions, usually the sigmoid.
- To optimize, you need gradient descent. For example, for \mathbf{W}_1

$$\mathbf{W}_1^{t+1} = \mathbf{W}_1^t - \alpha \nabla J(\mathbf{W}_1^t)$$

- But you need to do this for all $\mathbf{W}_1, \dots, \mathbf{W}_L$.
- And there are lots of sigmoid functions.
- Let us do the brute force.
- And this is back-propagation. (Really? Yes...)

Let us See an Example

- Let us look at two layers

$$J(\mathbf{W}_1, \mathbf{W}_2) = \left\| \underbrace{\sigma(\mathbf{W}_2^T \sigma(\mathbf{W}_1^T \mathbf{x})) - \mathbf{y}}_{\mathbf{a}_2} \right\|^2$$

- Let us go **backward**:

$$\frac{\partial J}{\partial \mathbf{W}_2} = \frac{\partial J}{\partial \mathbf{a}_2} \cdot \frac{\partial \mathbf{a}_2}{\partial \mathbf{W}_2}$$

- Now, what is \mathbf{a}_2 ?

$$\mathbf{a}_2 = \sigma(\underbrace{\mathbf{W}_2^T \sigma(\mathbf{W}_1^T \mathbf{x})}_{\mathbf{z}_2})$$

So let us compute:

$$\frac{\partial \mathbf{a}_2}{\partial \mathbf{W}_2} = \frac{\partial \mathbf{a}_2}{\partial \mathbf{z}_2} \cdot \frac{\partial \mathbf{z}_2}{\partial \mathbf{W}_2}.$$

Let us See an Example

$$J(\mathbf{W}_1, \mathbf{W}_2) = \|\sigma(\mathbf{W}_2^T \underbrace{\sigma(\mathbf{W}_1^T \mathbf{x})}_{\mathbf{a}_1}) - \mathbf{y}\|^2$$

- How about \mathbf{W}_1 ? Again, let us go **backward**:

$$\frac{\partial J}{\partial \mathbf{W}_1} = \frac{\partial J}{\partial \mathbf{a}_2} \cdot \frac{\partial \mathbf{a}_2}{\partial \mathbf{W}_1}$$

- But you can now repeat the calculation as follows (Let $\mathbf{z}_1 = \mathbf{W}_1^T \mathbf{x}$)


$$\begin{aligned}\frac{\partial \mathbf{a}_2}{\partial \mathbf{W}_1} &= \frac{\partial \mathbf{a}_2}{\partial \mathbf{a}_1} \frac{\partial \mathbf{a}_1}{\partial \mathbf{W}_1} \\ &= \frac{\partial \mathbf{a}_2}{\partial \mathbf{a}_1} \frac{\partial \mathbf{a}_1}{\partial \mathbf{z}_1} \frac{\partial \mathbf{z}_1}{\partial \mathbf{W}_1}\end{aligned}$$

- So it is just a very long sequence of chain rule.

Notations for Back Propagation

- The following notations are based on Nielsen's online book.
- The purpose of doing these is to write down a concise algorithm.

Weights:




w_{jk}^l is the weight from the k^{th} neuron in the $(l - 1)^{\text{th}}$ layer to the j^{th} neuron in the l^{th} layer

- w_{24}^3 : The 3rd layer
- w_{24}^3 : From 4-th neuron to 2-nd neuron

Notations for Back Propagation

Activation and Bias:



- a_1^3 : 3rd layer, 1st activation
- b_3^2 : 2nd layer, 3rd bias
- Here is the relationship. Think of $\sigma(\mathbf{w}^T \mathbf{x} + w_0)$:


$$a_j^\ell = \sigma \left(\sum_k w_{jk}^\ell a_k^{\ell-1} + b_j^\ell \right).$$

Understanding Back Propagation

- This is the main equation

$$a_j^\ell = \sigma \underbrace{\left(\sum_k w_{jk}^\ell a_k^{\ell-1} + b_j^\ell \right)}_{z_j^\ell}, \quad \text{or} \quad a_j^\ell = \sigma(z_j^\ell).$$

- a_j^ℓ : activation, z_j^ℓ : intermediate.




Loss

- The loss takes the form of

$$C = \sum_j (a_j^L - y_j)^2$$

- Think of two-class cross-entropy where each a^L is a 2-by-1 vector




$$C = (a_1^L - y_1)^2 + (a_2^L - y_2)^2$$

Error Term

- The error is defined as

$$\delta_j^\ell = \frac{\partial C}{\partial z_j^\ell}$$



- You can show that at the output,

$$\delta_j^L = \frac{\partial C}{\partial a_j^L} \frac{\partial a_j^L}{\partial z_j^L} = \frac{\partial C}{\partial a_j^L} \sigma'(z_j^L).$$

4 Fundamental Equations for Back Propagation

BP Equation 1: For the error in the output layer:

$$\delta_j^L = \frac{\partial C}{\partial a_j^L} \sigma'(z_j^L). \quad (\text{BP-1})$$

- First term: $\frac{\partial C}{\partial a_j^L}$ is rate of change w.r.t. a_j^L
- Second term: $\sigma'(z_j^L)$ = rate of change w.r.t. z_j^L .
- So it is just **chain rule**.
- Example: If $C = \frac{1}{2} \sum_j (y_j - a_j^L)^2$, then

$$\frac{\partial C}{\partial a_j^L} = (a_j^L - y_j)$$


- Matrix-vector form: $\boldsymbol{\delta}^L = \nabla_a C \odot \sigma'(\mathbf{z}^L)$

4 Fundamental Equations for Back Propagation

BP Equation 2: An equation for the error δ^ℓ in terms of the error in the next layer, $\delta^{\ell+1}$

$$\delta^\ell = ((\mathbf{w}^{\ell+1})^T \delta^{\ell+1}) \odot \sigma'(\mathbf{z}^\ell). \quad (\text{BP-2})$$

- You start with $\delta^{\ell+1}$. Take weighted average $\mathbf{w}^{\ell+1}$.



- (BP-1) and (BP-2) can help you determine error at any layer.

4 Fundamental Equations for Back Propagation

Equation 3: An equation for the rate of change of the cost with respect to any bias in the network.

$$\frac{\partial C}{\partial b_j^\ell} = \delta_j^\ell. \quad (\text{BP-3})$$

- Good news: We have already known δ_j^ℓ from Equation 1na dn 2.
- So computing $\frac{\partial C}{\partial b_j^\ell}$ is easy.

Equation 4: An equation for the rate of change of the cost with respect to any weight in the network.

$$\frac{\partial C}{\partial w_{jk}^\ell} = a_k^{\ell-1} \delta_j^\ell \quad (\text{BP-4})$$

- Again, everything on the right is known. So it is easy to compute.

Back Propagation Algorithm

- Below is a very concise summary of the BP algorithm


- 1. Input x :** Set the corresponding activation a^1 for the input layer.
- 2. Feedforward:** For each $l = 2, 3, \dots, L$ compute
$$z^l = w^l a^{l-1} + b^l \text{ and } a^l = \sigma(z^l).$$
- 3. Output error δ^L :** Compute the vector $\delta^L = \nabla_a C \odot \sigma'(z^L).$
- 4. Backpropagate the error:** For each $l = L-1, L-2, \dots, 2$ compute
$$\delta^l = ((w^{l+1})^T \delta^{l+1}) \odot \sigma'(z^l).$$
- 5. Output:** The gradient of the cost function is given by
$$\frac{\partial C}{\partial w_{jk}^l} = a_k^{l-1} \delta_j^l \text{ and } \frac{\partial C}{\partial b_j^l} = \delta_j^l.$$

Step 2: Feed Forward Step

- Let us take a closer look at Step 2
- The feed forward step computes the intermediate variables and the activations

$$z^\ell = (\mathbf{w}^\ell)^T \mathbf{a}^{\ell-1} + b^\ell$$

$$\mathbf{a}^\ell = \sigma(z^\ell).$$




$$a_j^l = \sigma(z_j^l) = \sigma((\mathbf{w}^l)^T \mathbf{a}^{l-1} + b_j^l)$$

Step 3: Output Error

- Let us take a closer look at Step 3
- The output error is given by (BP-1)


$$\delta^L = \nabla_a C \odot \sigma'(\mathbf{z}^L)$$



Step 4: Output Error

- Let us take a closer look at Step 4
- The error back propagation is given by (BP-2)

$$\delta^\ell = ((\mathbf{w}^{\ell+1})^T \delta^{\ell+1}) \odot \sigma'(\mathbf{z}^\ell).$$



Summary of Back Propagation

- There is no dark magic behind back propagation
- It is literally just **chain rule**
- You need to do this chain rule very systematically and **carefully**
- Then you can derive the back propagation steps
- Nielsen wrote in his book that

... How backpropagation could have been discovered in the first place? In fact, if you follow the approach I just sketched you will discover a proof of backpropagation... You make those simplifications, get a shorter proof, and write that out.... The result after a few iterations is the one we saw earlier, short but somewhat obscure...
- Most deep learning libraries have built-in back propagation steps.
- You don't have to implement it yourself, but you need to know what's behind it.

Reading List

- Michael Nielsen, Neural Networks and Deep Learning,
<http://neuralnetworksanddeeplearning.com/chap2.html>
 - Very well written. Easy to follow.
- Duda, Hart, Stork, Pattern Classification, Chapter 5
 - Classical treatment. Comprehensive. Readable.
- Bishop, Pattern Recognition and Machine Learning, Chapter 5
 - Somewhat Bayesian. Good for those who like statistics
- Stanford CS 231N, http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture4.pdf
 - Good numerical example.
- CMU <https://www.cs.cmu.edu/~mgormley/courses/10601-s17/slides/lecture20-backprop.pdf>
- Cornell <https://www.cs.cornell.edu/courses/cs5740/2016sp/resources/backprop.pdf>